

Electric Vehicle Data Analysis Project

August 16, 2025

1 *****Electric Vehicle Data Analysis*****

2 Import library

```
[2]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

3 Import dataset

```
[3]: df = pd.read_csv("Auta .csv")
```

3.0.1 Top Five data show

```
[4]: df.head()
```

```
[4]:
```

	Car full name	Make	Model	\
0	Audi e-tron 55 quattro	Audi	e-tron 55 quattro	
1	Audi e-tron 50 quattro	Audi	e-tron 50 quattro	
2	Audi e-tron S quattro	Audi	e-tron S quattro	
3	Audi e-tron Sportback 50 quattro	Audi	e-tron Sportback 50 quattro	
4	Audi e-tron Sportback 55 quattro	Audi	e-tron Sportback 55 quattro	

	Minimal price (gross) [PLN]	Engine power [KM]	Maximum torque [Nm]	\
0	345700	360	664	
1	308400	313	540	
2	414900	503	973	
3	319700	313	540	
4	357000	360	664	

	Type of brakes	Drive type	Battery capacity [kWh]	Range (WLTP) [km]	\
0	disc (front + rear)	4WD	95.0	438	
1	disc (front + rear)	4WD	71.0	340	
2	disc (front + rear)	4WD	95.0	364	
3	disc (front + rear)	4WD	71.0	346	

4	disc (front + rear)	4WD	95.0	447
---	---------------------	-----	------	-----

...	Permissable gross weight [kg]	Maximum load capacity [kg]	\
0	...	3130.0	640.0
1	...	3040.0	670.0
2	...	3130.0	565.0
3	...	3040.0	640.0
4	...	3130.0	670.0

	Number of seats	Number of doors	Tire size [in]	Maximum speed [kph]	\
0	5	5	19	200	
1	5	5	19	190	
2	5	5	20	210	
3	5	5	19	190	
4	5	5	19	200	

	Boot capacity (VDA) [l]	Acceleration 0-100 kph [s]	\
0	660.0	5.7	
1	660.0	6.8	
2	660.0	4.5	
3	615.0	6.8	
4	615.0	5.7	

	Maximum DC charging power [kW]	mean - Energy consumption [kWh/100 km]
0	150	24.45
1	150	23.80
2	150	27.55
3	150	23.30
4	150	23.85

[5 rows x 25 columns]

3.0.2 Change Column Name

```
[5]: # Column renaming mapping
col_map = {
    "Car full name": "car_full_name",
    "Make": "make",
    "Model": "model",
    "Minimal price (gross) [PLN]": "price_pln",
    "Engine power [KM]": "engine_power_km",
    "Maximum torque [Nm]": "max_torque_nm",
    "Type of brakes": "brake_type",
    "Drive type": "drive_type",
    "Battery capacity [kWh]": "battery_kwh",
    "Range (WLTP) [km]": "range_km",
    "Wheelbase [cm]": "wheelbase_cm",
}
```

```

    "Length [cm]": "length_cm",
    "Width [cm]": "width_cm",
    "Height [cm]": "height_cm",
    "Minimal empty weight [kg]": "empty_weight_kg",
    "Permissable gross weight [kg]": "gross_weight_kg",
    "Maximum load capacity [kg]": "max_load_kg",
    "Number of seats": "seats",
    "Number of doors": "doors",
    "Tire size [in]": "tire_in",
    "Maximum speed [kph]": "max_speed_kph",
    "Boot capacity (VDA) [l]": "boot_l",
    "Acceleration 0-100 kph [s]": "accel_0_100_s",
    "Maximum DC charging power [kW]": "dc_power_kw",
    "mean - Energy consumption [kWh/100 km]": "energy_kwh_per_100km"
}

```

```

# Apply renaming
df = df.rename(columns=col_map)

```

```
[6]: df.head(1)
```

```

[6]:      car_full_name  make      model  price_pln \
0  Audi e-tron 55 quattro  Audi  e-tron 55 quattro    345700

      engine_power_km  max_torque_nm      brake_type drive_type \
0              360          664  disc (front + rear)      4WD

      battery_kwh  range_km  ...  gross_weight_kg  max_load_kg  seats  doors \
0           95.0      438  ...          3130.0          640.0      5      5

      tire_in  max_speed_kph  boot_l  accel_0_100_s  dc_power_kw \
0          19           200  660.0           5.7          150

      energy_kwh_per_100km
0              24.45

[1 rows x 25 columns]

```

3.0.3 Check null value

```
[7]: df.isnull().sum()
```

```

[7]: car_full_name      0
     make              0
     model             0
     price_pln         0
     engine_power_km    0

```

```

max_torque_nm      0
brake_type         1
drive_type         0
battery_kwh        0
range_km           0
wheelbase_cm       0
length_cm          0
width_cm           0
height_cm          0
empty_weight_kg     0
gross_weight_kg     8
max_load_kg        8
seats              0
doors              0
tire_in            0
max_speed_kph      0
boot_l             1
accel_0_100_s       3
dc_power_kw        0
energy_kwh_per_100km 9
dtype: int64

```

```

[8]: # 1. Fill brake_type missing with mode
df["brake_type"] = df["brake_type"].fillna(df["brake_type"].mode()[0])

# 2. Fill gross_weight_kg missing with median
df["gross_weight_kg"] = df["gross_weight_kg"].fillna(df["gross_weight_kg"].
    ↪median())

# 3. Fill max_load_kg by calculation where possible, else median
mask = df["max_load_kg"].isna() & df["gross_weight_kg"].notna() &
    ↪df["empty_weight_kg"].notna()
df.loc[mask, "max_load_kg"] = df.loc[mask, "gross_weight_kg"] - df.loc[mask,
    ↪"empty_weight_kg"]
df["max_load_kg"] = df["max_load_kg"].fillna(df["max_load_kg"].median())

# 4. Fill boot_l with median
df["boot_l"] = df["boot_l"].fillna(df["boot_l"].median())

# 5. Fill accel_0_100_s with median per make
df["accel_0_100_s"] = df.groupby("make")["accel_0_100_s"].transform(lambda x: x.
    ↪fillna(x.median()))

# 6. Fill energy_kwh_per_100km by calculation if possible, else median
mask = df["energy_kwh_per_100km"].isna() & df["battery_kwh"].notna() &
    ↪df["range_km"].notna()

```

```
df.loc[mask, "energy_kwh_per_100km"] = (df.loc[mask, "battery_kwh"] / df.
↳loc[mask, "range_km"]) * 100
df["energy_kwh_per_100km"] = df["energy_kwh_per_100km"].
↳fillna(df["energy_kwh_per_100km"].median())
```

```
[9]: df.isnull().sum()
```

```
[9]: car_full_name      0
make                  0
model                 0
price_pln             0
engine_power_km       0
max_torque_nm         0
brake_type            0
drive_type            0
battery_kwh           0
range_km              0
wheelbase_cm          0
length_cm             0
width_cm              0
height_cm             0
empty_weight_kg       0
gross_weight_kg       0
max_load_kg           0
seats                 0
doors                 0
tire_in              0
max_speed_kph         0
boot_l               0
accel_0_100_s         0
dc_power_kw           0
energy_kwh_per_100km  0
dtype: int64
```

3.0.4 Find Information

```
[10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 53 entries, 0 to 52
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_full_name         53 non-null    object
1   make                  53 non-null    object
2   model                 53 non-null    object
3   price_pln             53 non-null    int64
4   engine_power_km       53 non-null    int64
```

```

5   max_torque_nm      53 non-null    int64
6   brake_type         53 non-null    object
7   drive_type         53 non-null    object
8   battery_kwh        53 non-null    float64
9   range_km           53 non-null    int64
10  wheelbase_cm       53 non-null    float64
11  length_cm          53 non-null    float64
12  width_cm           53 non-null    float64
13  height_cm          53 non-null    float64
14  empty_weight_kg    53 non-null    int64
15  gross_weight_kg    53 non-null    float64
16  max_load_kg        53 non-null    float64
17  seats              53 non-null    int64
18  doors              53 non-null    int64
19  tire_in            53 non-null    int64
20  max_speed_kph      53 non-null    int64
21  boot_l             53 non-null    float64
22  accel_0_100_s     53 non-null    float64
23  dc_power_kw        53 non-null    int64
24  energy_kwh_per_100km 53 non-null    float64
dtypes: float64(10), int64(10), object(5)
memory usage: 10.5+ KB

```

3.0.5 Describe

```
[11]: df.describe()
```

```

[11]:
count      price_pln  engine_power_km  max_torque_nm  battery_kwh  range_km  \
count      53.000000      53.000000      53.000000      53.000000      53.000000
mean    246158.509434      269.773585      460.037736      62.366038      376.905660
std    149187.485190      181.298589      261.647000      24.170913      118.817938
min      82050.000000      82.000000      160.000000      17.600000      148.000000
25%    142900.000000      136.000000      260.000000      40.000000      289.000000
50%    178400.000000      204.000000      362.000000      58.000000      364.000000
75%    339480.000000      372.000000      640.000000      80.000000      450.000000
max     794000.000000      772.000000     1140.000000     100.000000     652.000000

count      wheelbase_cm  length_cm  width_cm  height_cm  empty_weight_kg  \
count      53.000000      53.000000      53.000000      53.000000      53.000000
mean     273.581132     442.509434     186.241509     155.422642     1868.452830
std       22.740518      48.863280      14.280641      11.275358      470.880867
min       187.300000     269.500000     164.500000     137.800000     1035.000000
25%       258.800000     411.800000     178.800000     148.100000     1530.000000
50%       270.000000     447.000000     180.900000     155.600000     1685.000000
75%       290.000000     490.100000     193.500000     161.500000     2370.000000
max       327.500000     514.000000     255.800000     191.000000     2710.000000

count      gross_weight_kg  max_load_kg  seats  doors  tire_in  \
count      53.000000      53.000000      53.000000      53.000000      53.000000
mean     1868.452830      1868.452830      5.000000      4.000000      15.000000
std       470.880867      470.880867      1.000000      1.000000      1.000000
min       1035.000000      1035.000000      4.000000      3.000000      13.000000
25%       1530.000000      1530.000000      4.000000      3.000000      14.000000
50%       1685.000000      1685.000000      5.000000      4.000000      15.000000
75%       2370.000000      2370.000000      5.000000      4.000000      16.000000
max       2710.000000      2710.000000      5.000000      4.000000      17.000000

```

count	53.000000	53.000000	53.000000	53.000000	53.000000
mean	2263.207547	446.773585	4.905660	4.849057	17.679245
std	516.756614	264.158944	0.838133	0.455573	1.868500
min	1310.000000	-405.000000	2.000000	3.000000	14.000000
25%	1970.000000	417.000000	5.000000	5.000000	16.000000
50%	2119.000000	485.000000	5.000000	5.000000	17.000000
75%	2725.000000	565.000000	5.000000	5.000000	19.000000
max	3500.000000	1056.000000	8.000000	5.000000	21.000000

	max_speed_kph	boot_1	accel_0_100_s	dc_power_kw	\
count	53.000000	53.000000	53.000000	53.000000	
mean	178.169811	444.716981	7.332075	113.509434	
std	43.056196	178.458935	2.724940	57.166970	
min	123.000000	171.000000	2.500000	22.000000	
25%	150.000000	315.000000	5.100000	100.000000	
50%	160.000000	425.000000	7.600000	100.000000	
75%	200.000000	543.000000	9.000000	150.000000	
max	261.000000	870.000000	13.100000	270.000000	

	energy_kwh_per_100km
count	53.000000
mean	18.329014
std	4.357560
min	12.558140
25%	15.400000
50%	16.650000
75%	21.850000
max	28.200000

3.0.6 Find Duplicate

```
[12]: df.duplicated().any()
```

```
[12]: False
```

3.0.7 Last Five row

```
[13]: df.tail()
```

```
[13]:
```

	car_full_name	make	model	price_pln	\
48	Volkswagen ID.3 Pro S	Volkswagen	ID.3 Pro S	179990	
49	Volkswagen ID.4 1st	Volkswagen	ID.4 1st	202390	
50	Citroën ë-Spacetourer (M)	Citroën	ë-Spacetourer (M)	215400	
51	Mercedes-Benz EQV (long)	Mercedes-Benz	EQV (long)	339480	
52	Nissan e-NV200 evalia	Nissan	e-NV200 evalia	164328	

	engine_power_km	max_torque_nm	brake_type	drive_type	\
--	-----------------	---------------	------------	------------	---

48	204	310	disc (front) + drum (rear)	2WD (rear)
49	204	310	disc (front) + drum (rear)	2WD (rear)
50	136	260	disc (front + rear)	2WD (front)
51	204	362	disc (front + rear)	2WD (front)
52	109	254	disc (front + rear)	2WD (front)

	battery_kwh	range_km	...	gross_weight_kg	max_load_kg	seats	doors	\
48	77.0	549	...	2280.0	412.0	5	5	
49	77.0	500	...	2660.0	661.0	5	5	
50	50.0	230	...	2810.0	1056.0	8	5	
51	90.0	356	...	3500.0	865.0	6	5	
52	40.0	200	...	2250.0	658.0	5	5	

	tire_in	max_speed_kph	boot_l	accel_0_100_s	dc_power_kw	\
48	19	160	385.0	7.9	125	
49	20	160	543.0	8.5	125	
50	16	130	603.0	13.1	100	
51	17	160	425.0	5.1	110	
52	15	123	870.0	7.4	50	

	energy_kwh_per_100km
48	15.9
49	18.0
50	25.2
51	28.2
52	25.9

[5 rows x 25 columns]

3.0.8 Cleaning Process

```
[14]: import unicodecode
# 1. Fix encoding issues in text columns
for col in ["car_full_name", "make", "model"]:
    df[col] = (
        df[col]
        .astype(str) # ensure text
        .str.strip() # remove spaces
        .apply(lambda x: x.encode('latin1', errors='ignore').decode('utf-8',
errors='ignore'))
        .apply(unicodecode.unidecode) # remove accents if needed
    )
```

```
[15]: # Clean Drive Type
df["drive_type"] = df["drive_type"].str.strip().str.upper() # Remove spaces,
make uppercase
```



```
# Clean Brake Type
df["brake_type"] = df["brake_type"].str.strip().str.lower() # Remove spaces, ↵
↳make lowercase
df["brake_type"] = df["brake_type"].replace({
    "disc (front + rear)": "disc_all",
    "disc (front) + drum (rear)": "disc_drum"
})
```

```
[16]: df.head(3)
```

```
[16]:
```

	car_full_name	make	model	price_pln	\
0	Audi e-tron 55 quattro	Audi	e-tron 55 quattro	345700	
1	Audi e-tron 50 quattro	Audi	e-tron 50 quattro	308400	
2	Audi e-tron S quattro	Audi	e-tron S quattro	414900	

	engine_power_km	max_torque_nm	brake_type	drive_type	battery_kwh	\
0	360	664	disc_all	4WD	95.0	
1	313	540	disc_all	4WD	71.0	
2	503	973	disc_all	4WD	95.0	

	range_km	...	gross_weight_kg	max_load_kg	seats	doors	tire_in	\
0	438	...	3130.0	640.0	5	5	19	
1	340	...	3040.0	670.0	5	5	19	
2	364	...	3130.0	565.0	5	5	20	

	max_speed_kph	boot_l	accel_0_100_s	dc_power_kw	energy_kwh_per_100km
0	200	660.0	5.7	150	24.45
1	190	660.0	6.8	150	23.80
2	210	660.0	4.5	150	27.55

[3 rows x 25 columns]

```
[17]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 53 entries, 0 to 52
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_full_name          53 non-null    object
1   make                   53 non-null    object
2   model                   53 non-null    object
3   price_pln              53 non-null    int64
4   engine_power_km        53 non-null    int64
5   max_torque_nm          53 non-null    int64
6   brake_type             53 non-null    object
7   drive_type             53 non-null    object
8   battery_kwh            53 non-null    float64
```

```

9   range_km          53 non-null    int64
10  wheelbase_cm      53 non-null    float64
11  length_cm         53 non-null    float64
12  width_cm          53 non-null    float64
13  height_cm         53 non-null    float64
14  empty_weight_kg   53 non-null    int64
15  gross_weight_kg   53 non-null    float64
16  max_load_kg       53 non-null    float64
17  seats             53 non-null    int64
18  doors             53 non-null    int64
19  tire_in           53 non-null    int64
20  max_speed_kph     53 non-null    int64
21  boot_l            53 non-null    float64
22  accel_0_100_s     53 non-null    float64
23  dc_power_kw       53 non-null    int64
24  energy_kwh_per_100km 53 non-null    float64

```

dtypes: float64(10), int64(10), object(5)

memory usage: 10.5+ KB

3.0.9 Task 1: A customer has a budget of 350,000 PLN and wants an EV with a minimum range of 400 km.

a) Your task is to filter out EVs that meet these criteria.

```

[18]: # Filter EVs
filtered_df = df[
    (df["price_pln"] <= 350000) &
    (df["range_km"] >= 400)
]

# Show only key columns
print(filtered_df[["car_full_name", "make", "model", "price_pln", "range_km",
    ↪ "battery_kwh"]])

```

	car_full_name	make \
0	Audi e-tron 55 quattro	Audi
8	BMW iX3	BMW
15	Hyundai Kona electric 64kWh	Hyundai
18	Kia e-Niro 64kWh	Kia
20	Kia e-Soul 64kWh	Kia
22	Mercedes-Benz EQC	Mercedes-Benz
39	Tesla Model 3 Standard Range Plus	Tesla
40	Tesla Model 3 Long Range	Tesla
41	Tesla Model 3 Performance	Tesla
47	Volkswagen ID.3 Pro Performance	Volkswagen
48	Volkswagen ID.3 Pro S	Volkswagen
49	Volkswagen ID.4 1st	Volkswagen

	model	price_pln	range_km	battery_kwh
--	-------	-----------	----------	-------------

0	e-tron 55 quattro	345700	438	95.0
8	iX3	282900	460	80.0
15	Kona electric 64kWh	178400	449	64.0
18	e-Niro 64kWh	167990	455	64.0
20	e-Soul 64kWh	160990	452	64.0
22	EQC	334700	414	80.0
39	Model 3 Standard Range Plus	195490	430	54.0
40	Model 3 Long Range	235490	580	75.0
41	Model 3 Performance	260490	567	75.0
47	ID.3 Pro Performance	155890	425	58.0
48	ID.3 Pro S	179990	549	77.0
49	ID.4 1st	202390	500	77.0

3.0.10 b) Group them by the manufacturer (Make).

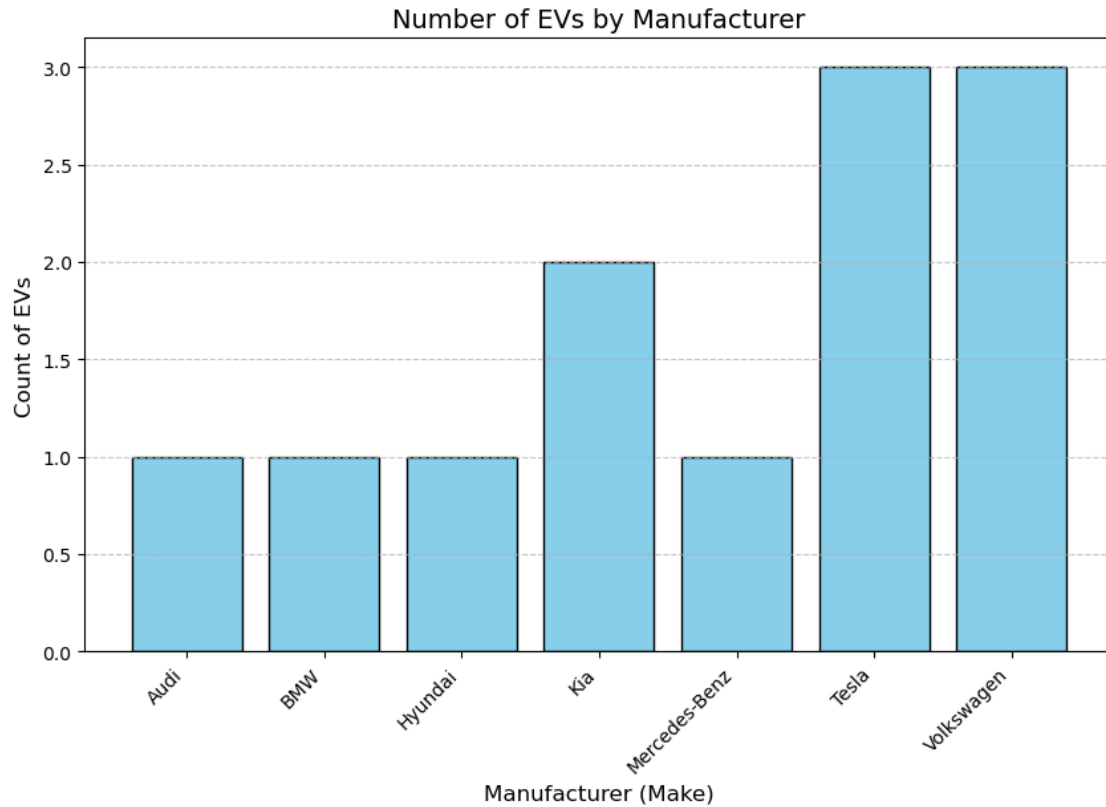
```
[19]: grouped_by_make = filtered_df.groupby("make").size().reset_index(name="count")
      print(grouped_by_make)
```

	make	count
0	Audi	1
1	BMW	1
2	Hyundai	1
3	Kia	2
4	Mercedes-Benz	1
5	Tesla	3
6	Volkswagen	3

```
[20]: # Bar plot of EV count per manufacturer
plt.figure(figsize=(10, 6))
plt.bar(grouped_by_make["make"], grouped_by_make["count"], color="skyblue",
        edgecolor="black")

plt.title("Number of EVs by Manufacturer", fontsize=14)
plt.xlabel("Manufacturer (Make)", fontsize=12)
plt.ylabel("Count of EVs", fontsize=12)
plt.xticks(rotation=45, ha="right")
plt.grid(axis="y", linestyle="--", alpha=0.7)

plt.show()
```



3.0.11 c) Calculate the average battery capacity for each manufacturer.

```
[21]: # Calculate average battery capacity by manufacturer
avg_battery_capacity = df.groupby("make")["battery_kwh"].mean().reset_index()

# Rename column for clarity
avg_battery_capacity.columns = ["make", "avg_battery_kwh"]

print(avg_battery_capacity)
```

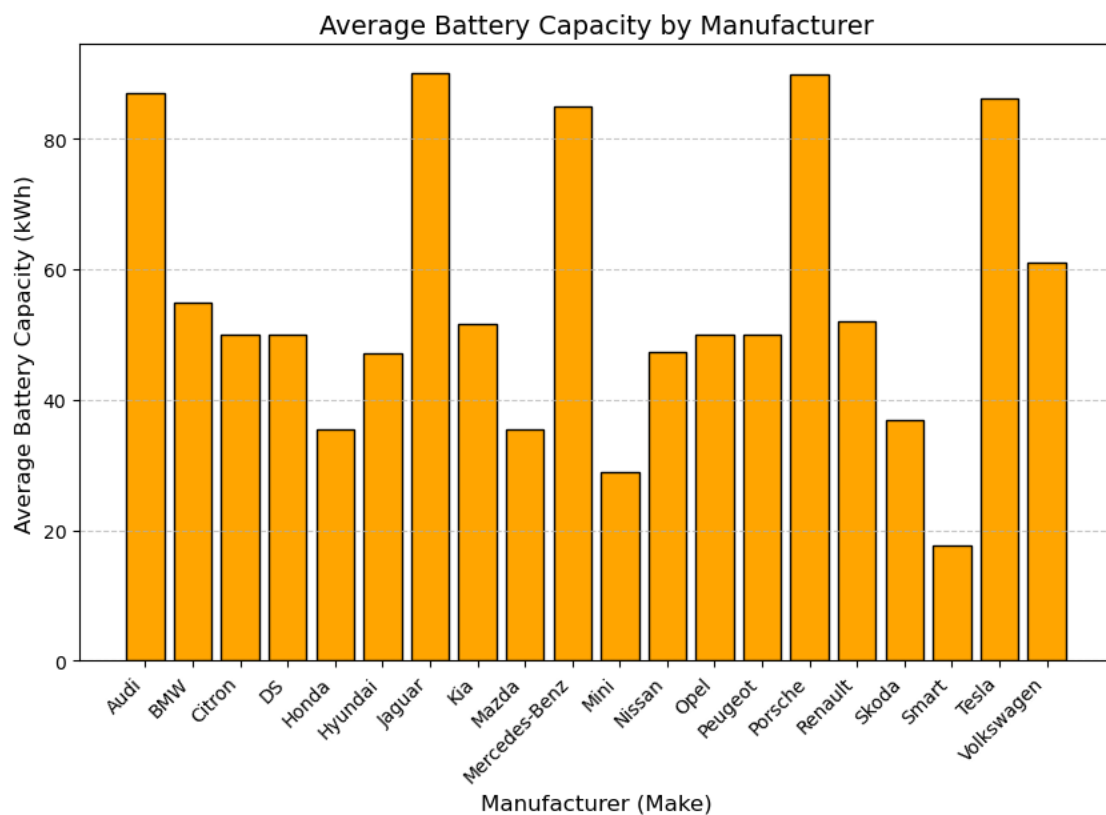
	make	avg_battery_kwh
0	Audi	87.000000
1	BMW	54.800000
2	Citron	50.000000
3	DS	50.000000
4	Honda	35.500000
5	Hyundai	47.166667
6	Jaguar	90.000000
7	Kia	51.600000
8	Mazda	35.500000
9	Mercedes-Benz	85.000000
10	Mini	28.900000

11	Nissan	47.333333
12	Opel	50.000000
13	Peugeot	50.000000
14	Porsche	89.850000
15	Renault	52.000000
16	Skoda	36.800000
17	Smart	17.600000
18	Tesla	86.285714
19	Volkswagen	61.075000

```
[22]: # Plot average battery capacity by manufacturer
plt.figure(figsize=(10, 6))
plt.bar(avg_battery_capacity["make"], avg_battery_capacity["avg_battery_kwh"],
        color="orange", edgecolor="black")

plt.title("Average Battery Capacity by Manufacturer", fontsize=14)
plt.xlabel("Manufacturer (Make)", fontsize=12)
plt.ylabel("Average Battery Capacity (kWh)", fontsize=12)
plt.xticks(rotation=45, ha="right")
plt.grid(axis="y", linestyle="--", alpha=0.7)

plt.show()
```



3.0.12 Task 2: You suspect some EVs have unusually high or low energy consumption. Find the outliers in the mean - Energy consumption [kWh/100 km] column

```
[23]: # Calculate Q1 and Q3
Q1 = df["energy_kwh_per_100km"].quantile(0.25)
Q3 = df["energy_kwh_per_100km"].quantile(0.75)
IQR = Q3 - Q1

# Define lower and upper bounds
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Filter outliers
outliers = df[(df["energy_kwh_per_100km"] < lower_bound) |
               (df["energy_kwh_per_100km"] > upper_bound)]

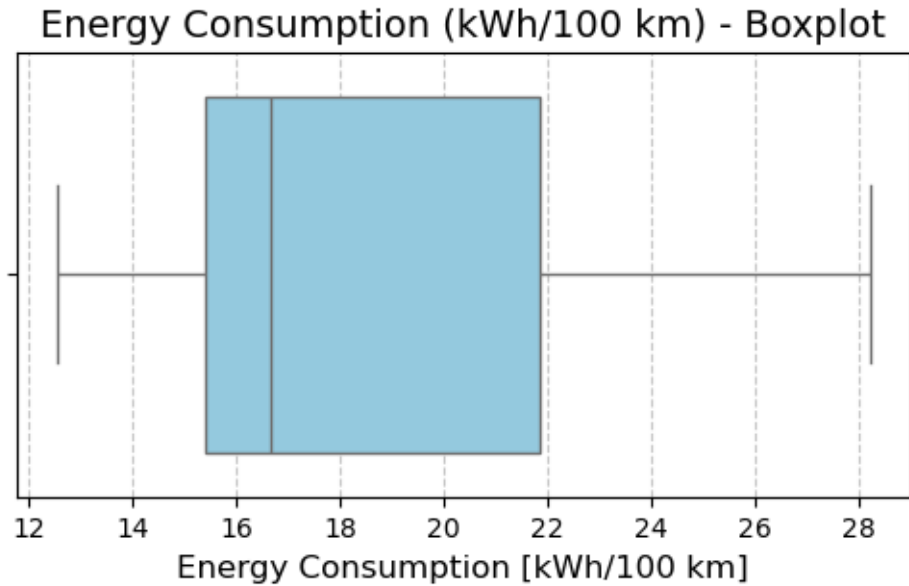
print("Outliers in energy consumption:")
print(outliers[["car_full_name", "make", "energy_kwh_per_100km"]])
if outliers.empty:
    print("No outliers detected in energy consumption using IQR method.")
```

```
Outliers in energy consumption:
Empty DataFrame
Columns: [car_full_name, make, energy_kwh_per_100km]
Index: []
No outliers detected in energy consumption using IQR method.
```

```
[24]: plt.figure(figsize=(6, 3))
sns.boxplot(x=df["energy_kwh_per_100km"], color="skyblue")

plt.title("Energy Consumption (kWh/100 km) - Boxplot", fontsize=14)
plt.xlabel("Energy Consumption [kWh/100 km]", fontsize=12)
plt.grid(axis="x", linestyle="--", alpha=0.7)

plt.show()
```



3.0.13 Show the all columns

```
[25]: df.columns
```

```
[25]: Index(['car_full_name', 'make', 'model', 'price_pln', 'engine_power_kmh',
            'max_torque_nm', 'brake_type', 'drive_type', 'battery_kwh', 'range_km',
            'wheelbase_cm', 'length_cm', 'width_cm', 'height_cm', 'empty_weight_kg',
            'gross_weight_kg', 'max_load_kg', 'seats', 'doors', 'tire_in',
            'max_speed_kph', 'boot_l', 'accel_0_100_s', 'dc_power_kw',
            'energy_kwh_per_100km'],
            dtype='object')
```

3.0.14 Task 3: Your manager wants to know if there's a strong relationship between battery capacity and range.

a) Create a suitable plot to visualize.


```
[26]: # Scatter plot for Battery Capacity vs Range
plt.figure(figsize=(8, 6))
plt.scatter(df["battery_kwh"], df["range_km"],
            color="royalblue", alpha=0.7, edgecolor="black")

# Titles and labels
plt.title("Battery Capacity vs Range", fontsize=14, fontweight="bold")
plt.xlabel("Battery Capacity (kWh)", fontsize=12)
plt.ylabel("Range (km)", fontsize=12)

# Grid for better readability
```

```
plt.grid(True, linestyle="--", alpha=0.6)

plt.show()
```



3.0.15 b) Highlight any insights.

Insights Positive correlation – Cars with higher battery capacity generally have a higher range, indicating that battery size is a strong driver of range. Diminishing returns – The relationship isn't perfectly linear; after around 80–100 kWh, increases in capacity yield smaller gains in range, possibly due to vehicle weight and efficiency limits.

Outliers – Some EVs achieve high range with a relatively smaller battery (likely due to better aerodynamics and energy efficiency). Others have large batteries but lower-than-expected range (possibly heavy SUVs or performance-oriented EVs). **Manufacturer effect** – Certain brands consistently position above the trend line, meaning they are more efficient at converting battery capacity into range.

3.0.16 Task 4: Build an EV recommendation class. The class should allow users to input their budget, desired range, and battery capacity. The class should then return the top three EVs matching their criteria

```
[27]: class EVRecommender:
    def __init__(self, df):
        self.df = df

    def recommend(self):
        # Show min-max values from dataset
        print(f"\nPrice range in dataset: {self.df['price_pln'].min()} - {self.
↪df['price_pln'].max()} PLN")
        print(f"Range (km) in dataset: {self.df['range_km'].min()} - {self.
↪df['range_km'].max()} km")
        print(f"Battery capacity in dataset: {self.df['battery_kwh'].min()} -
↪{self.df['battery_kwh'].max()} kWh\n")

        # Get user inputs
        budget = float(input("Enter your budget (PLN): "))
        min_range = float(input("Enter minimum desired range (km): "))
        min_battery = float(input("Enter minimum battery capacity (kWh): "))

        # Filter dataset based on criteria
        filtered = self.df[
            (self.df["price_pln"] <= budget) &
            (self.df["range_km"] >= min_range) &
            (self.df["battery_kwh"] >= min_battery)
        ]

        # Sort by range (highest first) and price (lowest first)
        filtered = filtered.sort_values(by=["range_km", "price_pln"],
↪ascending=[False, True])

        # Show results
        if filtered.empty:
            print("\nNo matching EVs found for your criteria.")
        else:
            print("\nTop 3 matching EVs:")
            print(filtered[["car_full_name", "make", "price_pln", "range_km",
↪"battery_kwh"]].head(3))

# Example usage:
#recommender = EVRecommender(df)
# recommender.recommend()
```

3.0.17 Task 5: Inferential Statistics – Hypothesis Testing: Test whether there is a significant difference in the average Engine power [KM] of vehicles manufactured by two leading manufacturers i.e. Tesla and Audi.

```
[28]: from scipy.stats import ttest_ind
import pandas as pd

# Load dataset
df = pd.read_csv("my_file.csv")

# Filter data for Tesla and Audi
tesla_power = df[df["make"] == "Tesla"]["engine_power_km"]
audi_power = df[df["make"] == "Audi"]["engine_power_km"]

# Perform Welch's t-test (does not assume equal variance)
t_stat, p_value = ttest_ind(tesla_power, audi_power, equal_var=False)

print("T-statistic:", round(t_stat, 3))
print("P-value:", round(p_value, 5))

# Interpretation
alpha = 0.05
if p_value < alpha:
    print("\n Conclusion: There IS a statistically significant difference in_
    ↪average engine power between Tesla and Audi.")
else:
    print("\n Conclusion: There is NO statistically significant difference in_
    ↪average engine power between Tesla and Audi.")
```

T-statistic: 1.794

P-value: 0.10684

Conclusion: There is NO statistically significant difference in average engine power between Tesla and Audi.

3.0.18 What insights can you draw from the test results?

3.0.19 Recommendations and Conclusion: Provide actionable insights based on your analysis.

Insight We compared the engine power of Tesla and Audi electric vehicles (EVs) using a t-test to see if one brand has consistently stronger engines. The T-statistic is 1.794, and the P-value is 0.10684. Since the P-value is greater than 0.05, the difference in engine power between Tesla and Audi is likely due to chance rather than a real difference.

Conclusion There is no statistically significant difference in average engine power between Tesla and Audi EVs. This means that, on average, both brands perform similarly in terms of engine power.

Recommendation Since there's no clear winner in engine power, customers should consider other

factors when choosing between Tesla and Audi EVs, such as price, driving range, charging speed, features, and brand preference, rather than focusing solely on engine power.

```
[30]: df.to_csv("Final.csv", index= False)
```

```
[ ]:
```